

sklearn.metrics.log_loss

```
sklearn.metrics.log_loss(y_true, y_pred, eps=1e-15, normalize=True, sample_weight=None)
```

[\[source\]](#)

Log loss, aka logistic loss or cross-entropy loss.

This is the loss function used in (multinomial) logistic regression and extensions of it such as neural networks, defined as the negative log-likelihood of the true labels given a probabilistic classifier's predictions. For a single sample with true label y_t in $\{0, 1\}$ and estimated probability y_p that $y_t = 1$, the log loss is

$$-\log P(y_t|y_p) = -(y_t \log(y_p) + (1 - y_t) \log(1 - y_p))$$

Parameters: **y_true** : array-like or label indicator matrix

Ground truth (correct) labels for n_{samples} samples.

y_pred : array-like of float, shape = (n_{samples} , n_{classes})

Predicted probabilities, as returned by a classifier's `predict_proba` method.

eps : float

Log loss is undefined for $p=0$ or $p=1$, so probabilities are clipped to $\max(\text{eps}, \min(1 - \text{eps}, p))$.

normalize : bool, optional (default=True)

If true, return the mean loss per sample. Otherwise, return the sum of the per-sample losses.

sample_weight : array-like of shape = [n_{samples}], optional

Sample weights.

Returns: **loss** : float

Notes

The logarithm used is the natural logarithm (base-e).

References


C.M. Bishop (2006). Pattern Recognition and Machine Learning. Springer, p. 209.

Examples

```
>>> log_loss(["spam", "ham", "ham", "spam"],
...          [[.1, .9], [.9, .1], [.8, .2], [.35, .65]])
0.21616...
```

[>>>](#)

Examples using sklearn.metrics.log_loss



Probability Calibration
for 3-class classification

